

Assignment_4

September 29, 2023

1 Model building on Employee Attrition Dataset

Anand Misra 21BAI1105

1.1 Importing the Libraries

```
[82]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
```

1.2 Importing the Dataset

```
[83]: data = pd.read_csv('Employee-Attrition.csv')
```

1.3 Checking for Null Values

```
[84]: null_counts = data.isnull().sum()
null_counts
```

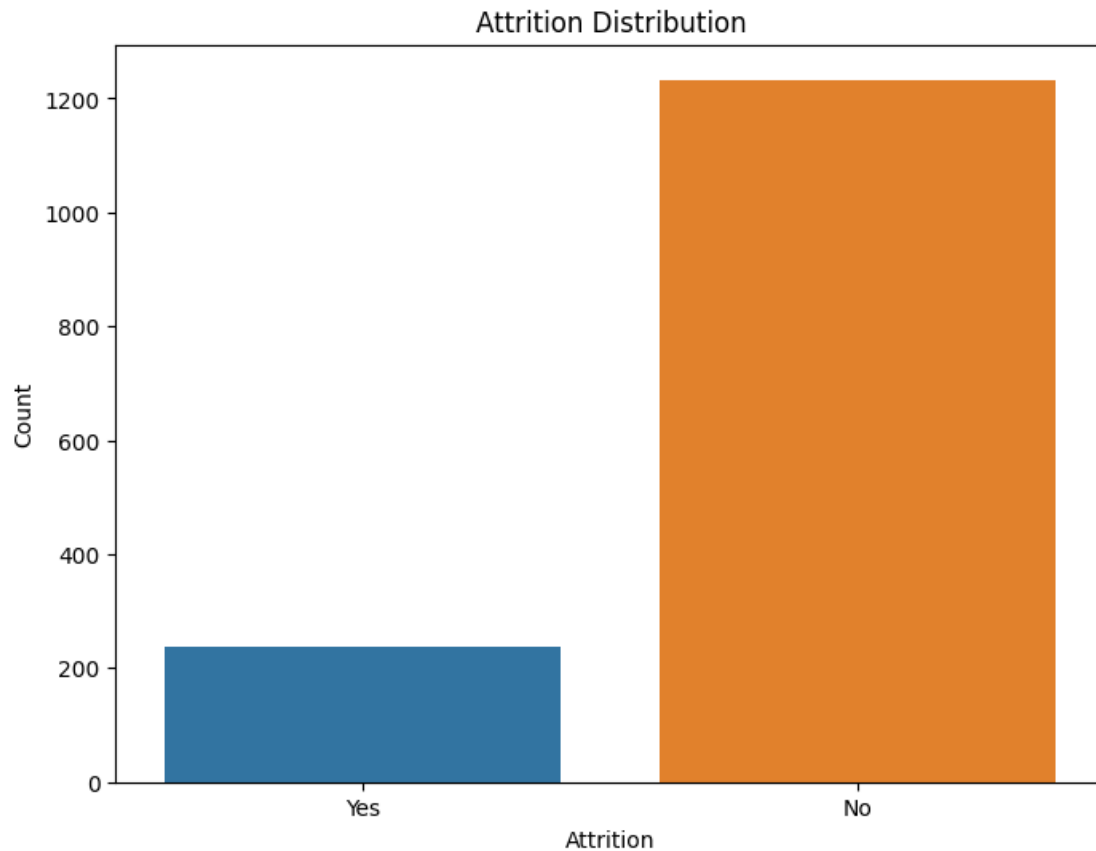
```
[84]: Age                                0
Attrition                              0
BusinessTravel                         0
DailyRate                             0
Department                             0
DistanceFromHome                       0
Education                              0
EducationField                          0
EmployeeCount                          0
EmployeeNumber                         0
EnvironmentSatisfaction                0
Gender                                 0
HourlyRate                             0
JobInvolvement                         0
JobLevel                               0
JobRole                                0
```

JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0
PercentSalaryHike	0
PerformanceRating	0
RelationshipSatisfaction	0
StandardHours	0
StockOptionLevel	0
TotalWorkingYears	0
TrainingTimesLastYear	0
WorkLifeBalance	0
YearsAtCompany	0
YearsInCurrentRole	0
YearsSinceLastPromotion	0
YearsWithCurrManager	0


dtype: int64

1.4 Data Visualization

```
[85]: plt.figure(figsize=(8, 6))
sns.countplot(data=data, x='Attrition')
plt.title('Attrition Distribution')
plt.xlabel('Attrition')
plt.ylabel('Count')
plt.show()
```



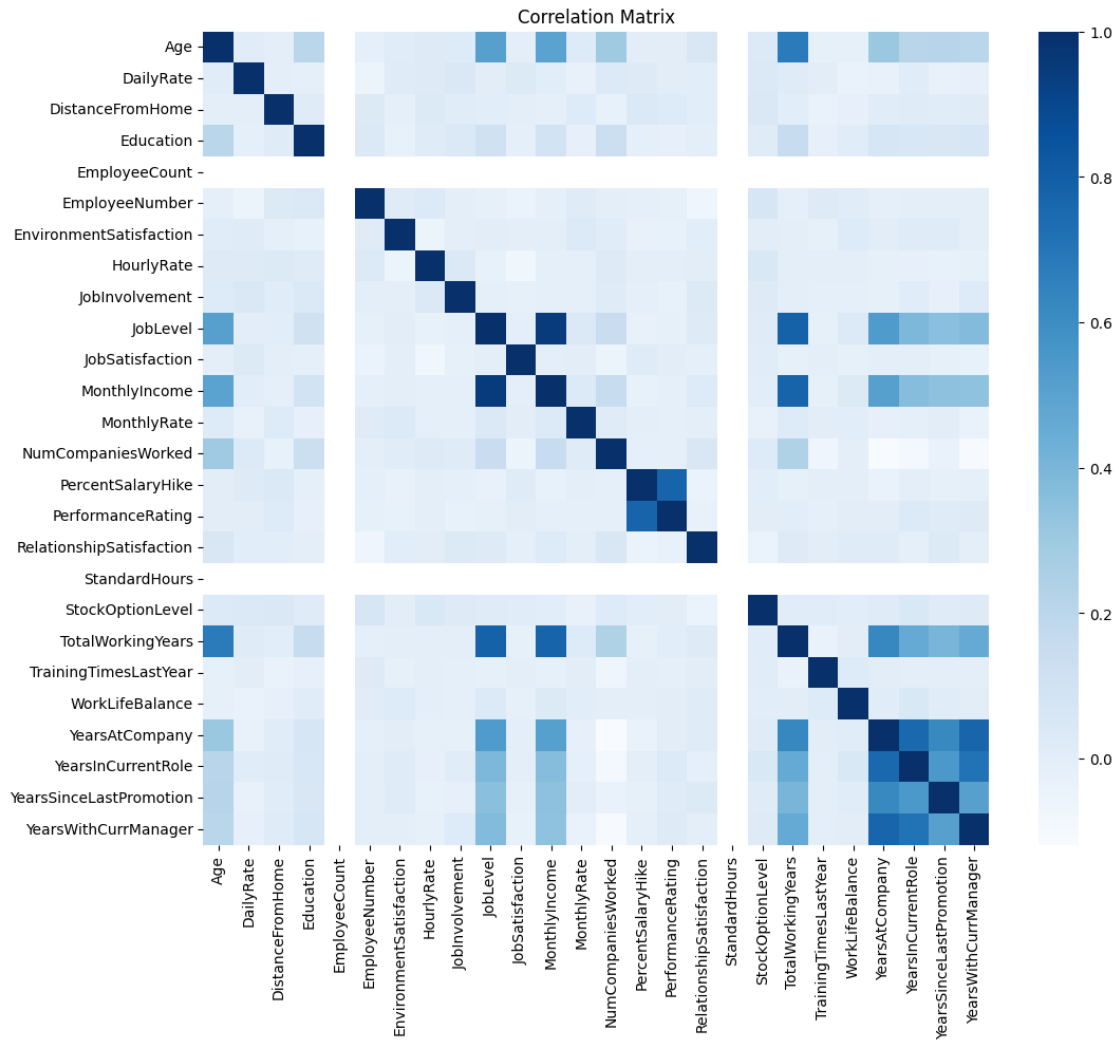
```
[86]: attrition_counts = data['Attrition'].value_counts()
labels = attrition_counts.index
sizes = attrition_counts.values
plt.figure(figsize=(6, 6))
plt.pie(sizes, labels=labels, autopct='%1.1f%%', startangle=90)
plt.title('Attrition Distribution')
plt.axis('equal')
plt.show()
```



```
[87]: correlation_matrix = data.corr()  
plt.figure(figsize=(12, 10))  
sns.heatmap(correlation_matrix, cmap='Blues')  
plt.title('Correlation Matrix')  
plt.show()
```

<ipython-input-87-9b4c1a72ae78>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = data.corr()
```



1.5 Removing columns that are not useful or not related

```
[88]: data = data.drop(['BusinessTravel', 'EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours', 'PerformanceRating'], axis=1)
```

1.6 Encoding the required attributes

```
[89]: data.head()
```

```
[89]:
```

	Age	Attrition	DailyRate	Department	DistanceFromHome	\
0	41	Yes	1102	Sales		1
1	49	No	279	Research & Development		8
2	37	Yes	1373	Research & Development		2
3	33	No	1392	Research & Development		3

4	27	No	591	Research & Development	2
---	----	----	-----	------------------------	---

	Education	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	...	\
0	2	Life Sciences	2	Female	94	...	
1	1	Life Sciences	3	Male	61	...	
2	2	Other	4	Male	92	...	
3	4	Life Sciences	4	Female	56	...	
4	1	Medical	1	Male	40	...	

	PercentSalaryHike	RelationshipSatisfaction	StockOptionLevel	\
0	11	1	0	
1	23	4	1	
2	15	2	0	
3	11	3	0	
4	12	4	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 29 columns]

```
[90]: label_encoders = {}
categorical_columns = data.select_dtypes(include=['object']).columns
for col in categorical_columns:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])
    label_encoders[col] = le
```

```
[91]: data.head()
```

	Age	Attrition	DailyRate	Department	DistanceFromHome	Education	\
0	41	1	1102	2	1	2	
1	49	0	279	1	8	1	
2	37	1	1373	1	2	2	
3	33	0	1392	1	3	4	
4	27	0	591	1	2	1	

	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	...	\
0	1	2	0	94	...	
1	1	3	1	61	...	
2	4	4	1	92	...	
3	1	4	0	56	...	
4	3	1	1	40	...	

	PercentSalaryHike	RelationshipSatisfaction	StockOptionLevel	\
0	11	1	0	
1	23	4	1	
2	15	2	0	
3	11	3	0	
4	12	4	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 29 columns]

1.7 Splitting Dependent and Independent variables

```
[92]: X = data.drop('Attrition', axis=1)
      y = data['Attrition']
```

```
[93]: X.head()
```

```
[93]:
```

	Age	DailyRate	Department	DistanceFromHome	Education	EducationField	\
0	41	1102	2	1	2	1	
1	49	279	1	8	1	1	
2	37	1373	1	2	2	4	
3	33	1392	1	3	4	1	
4	27	591	1	2	1	3	

	EnvironmentSatisfaction	Gender	HourlyRate	JobInvolvement	...	\
0	2	0	94	3	...	

1	3	1	61	2 ...
2	4	1	92	2 ...
3	4	0	56	3 ...
4	1	1	40	3 ...

	PercentSalaryHike	RelationshipSatisfaction	StockOptionLevel	\
0	11	1	0	
1	23	4	1	
2	15	2	0	
3	11	3	0	
4	12	4	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 28 columns]

```
[94]: y.head()
```

```
[94]: 0    1
      1    0
      2    1
      3    0
      4    0
      Name: Attrition, dtype: int64
```

1.8 Splitting Data into Train and Test

```
[95]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```


1.9 Logistic Regression

```
[96]: from sklearn.linear_model import LogisticRegression
```

```
logistic_model = LogisticRegression()  
logistic_model.fit(X_train, y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[96]: LogisticRegression()
```

1.10 Decision Tree

```
[97]: from sklearn.tree import DecisionTreeClassifier
```

```
decision_tree_model = DecisionTreeClassifier(random_state=42)  
decision_tree_model.fit(X_train, y_train)
```

```
[97]: DecisionTreeClassifier(random_state=42)
```

1.11 Performance Metrics

```
[98]: from sklearn.metrics import accuracy_score, classification_report,  
      ↪ confusion_matrix
```

1.11.1 Performance of logistic regression

```
[99]: logistic_predictions = logistic_model.predict(X_test)  
logistic_accuracy = accuracy_score(y_test, logistic_predictions)  
logistic_classification_report = classification_report(y_test,  
      ↪ logistic_predictions)  
logistic_confusion_matrix = confusion_matrix(y_test, logistic_predictions)
```

```
[100]: print("Logistic Regression Accuracy:", logistic_accuracy)  
print("\nLogistic Regression Classification Report:\n",  
      ↪ logistic_classification_report)  
print("\nLogistic Regression Confusion Matrix:\n", logistic_confusion_matrix)
```

Logistic Regression Accuracy: 0.8673469387755102

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.87	1.00	0.93	255
1	0.50	0.03	0.05	39
accuracy			0.87	294
macro avg	0.68	0.51	0.49	294
weighted avg	0.82	0.87	0.81	294

Logistic Regression Confusion Matrix:

```
[[254  1]
 [ 38  1]]
```

1.11.2 Performance of decision tree

```
[101]: decision_tree_predictions = decision_tree_model.predict(X_test)
decision_tree_accuracy = accuracy_score(y_test, decision_tree_predictions)
decision_tree_classification_report = classification_report(y_test,
    ↳decision_tree_predictions)
decision_tree_confusion_matrix = confusion_matrix(y_test,
    ↳decision_tree_predictions)
```

```
[102]: print("\nDecision Tree Accuracy:", decision_tree_accuracy)
print("\nDecision Tree Classification Report:\n",
    ↳decision_tree_classification_report)
print("\nDecision Tree Confusion Matrix:\n", decision_tree_confusion_matrix)
```

Decision Tree Accuracy: 0.7891156462585034

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.88	0.87	0.88	255
1	0.22	0.23	0.23	39
accuracy			0.79	294
macro avg	0.55	0.55	0.55	294
weighted avg	0.79	0.79	0.79	294

Decision Tree Confusion Matrix:

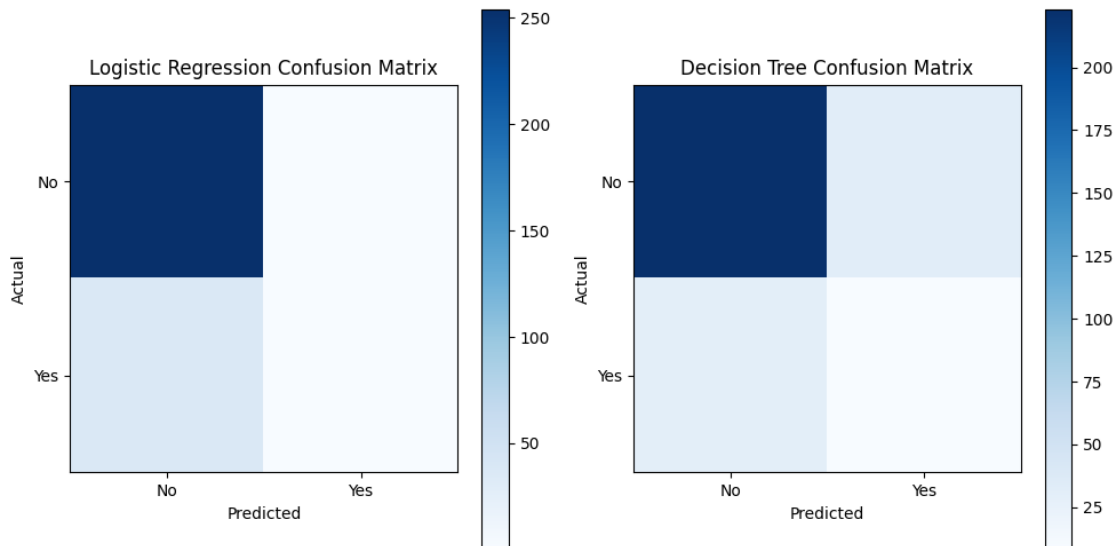
```
[[223  32]
 [ 30   9]]
```

1.11.3 Visualization of performance metrics

```
[103]: plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.title("Logistic Regression Confusion Matrix")
plt.imshow(logistic_confusion_matrix, cmap='Blues', interpolation='nearest')
plt.colorbar()
plt.xticks([0, 1], ['No', 'Yes'])
plt.yticks([0, 1], ['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')

plt.subplot(1, 2, 2)
plt.title("Decision Tree Confusion Matrix")
plt.imshow(decision_tree_confusion_matrix, cmap='Blues',
           interpolation='nearest')
plt.colorbar()
plt.xticks([0, 1], ['No', 'Yes'])
plt.yticks([0, 1], ['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')

plt.tight_layout()
plt.show()
```



```
[ ]:
```